

A Reproducible Local-Ratio and SCC-Refinement Framework for Weighted Ordering in Directed Graphs

Abstract

Weighted directed graphs model precedence conflicts, cyclic dependencies, and pairwise ordering data in engineering and decision systems. This paper studies the minimum weighted feedback arc set problem on nonnegative weighted directed graphs, equivalently the problem of finding an ordering with minimum total backward edge weight. We present a reproducible heuristic framework that combines an engineered local-ratio cycle-reduction method with topological add-back, a weighted seed construction, and incumbent-protected refinement over strongly connected components. The refinement improves local cyclic substructures while preserving a non-worsening invariant relative to the best seed solution.

The computational study combines sparse public graph benchmarks, an ablation study, exact validation on small instances, external baseline comparisons, and a dense linear-ordering transfer test. On small nonnegative instances where exact optimization is feasible, the proposed refinement is optimal on 56 of 57 standard instances, with a mean relative gap of about 0.0006 percent. On 97 standard nonnegative sparse benchmark instances, it attains the best observed backward weight on 96 instances and outperforms all tested baselines, including the strongest external heuristic. Dense instances from the Linear Ordering Problem Library (LOLIB) define an important scope boundary. A tournament-native external method is stronger on that benchmark, while the proposed framework remains a general weighted-digraph heuristic rather than a dense-native linear-ordering solver.

Keywords: Minimum weighted feedback arc set, weighted ordering, local-ratio algorithm, strongly connected components, incumbent-protected refinement, linear ordering problem, computational heuristics

1. Introduction

Many engineering and decision systems require a linear order to be extracted from directed, weighted, and internally inconsistent information. A precedence graph may encode that task u should precede task v , that a circuit signal should be routed before a dependent component is evaluated, that a package should be installed before another package can be configured, or that an alternative is preferred to another alternative in paired-comparison data. In practice these directed relations are often cyclic. Cycles can arise from noisy measurements, conflicting expert judgments, reciprocal dependencies, approximate model reductions, or competing design objectives. A useful ordering method must therefore decide which directed constraints to respect and which constraints to place backward in the final order. When arcs carry weights, the backward constraints are not interchangeable: violating a high-weight precedence relation should be more costly than violating a weak or noisy relation. This leads naturally to the minimum weighted feedback arc set (MWFAS) problem, a weighted ordering problem on directed graphs with direct relevance to cyclic-dependency resolution, ranking, scheduling, and other computational decision workflows in industrial engineering.

Informally, the minimum weighted feedback arc set problem asks for a minimum-weight set of arcs whose removal makes a directed graph acyclic. Equivalently, it asks for a vertex ordering that minimizes the total weight of arcs directed backward with respect to that order. The problem is computationally challenging, and exact methods are practical only for small instances or special structures. Naive score heuristics can be very fast, but they may ignore local cyclic structure; random multistart heuristics can provide a broad baseline, but they do not exploit graph decomposition; and specialized dense-ordering solvers may perform well on complete pairwise-comparison instances while being less directly aligned with sparse, irregular directed graphs. The computational question addressed here is not whether one heuristic dominates every possible ordering instance. Instead, it is whether a carefully engineered and reproducible framework can deliver strong performance on nonnegative sparse weighted digraphs while making its scope limits explicit.

The local-ratio perspective is an important starting point for feedback arc set heuristics and approximation-oriented reasoning. We do not claim local-ratio as new. Broader local-ratio foundations appear in the work of Bar-Yehuda and coauthors Bar-Yehuda et al. (1998), and directed feedback problems were treated algorithmically by Demetrescu and Finocchi Demetrescu

and Finocchi (2003). The contribution in this paper is algorithm-engineering rather than a new approximation theorem: we study a reproducible local-ratio cycle-reduction implementation that records cycle-removal decisions, reconstructs an acyclic topological order, and then performs a topological add-back phase to recover arcs whenever doing so does not reintroduce cyclicity. This method is referred to as local-ratio topological add-back (LR-TA) after the components are defined in Section 4. The emphasis is on deterministic data handling, measurable component effects, and integration with a refinement stage that is evaluated against exact and external baselines.

The integrated framework combines three components. First, the local-ratio topological add-back procedure produces a feasible ordering from weighted cycle reductions and a reconstruction step. Second, a weighted seed construction provides a complementary initial ordering; in this manuscript it is treated as a seed and baseline rather than as a primary novelty claim.

Third, incumbent-protected refinement searches within strongly connected components (SCCs). It updates the incumbent only when a candidate ordering has no larger backward weight. We refer to this refinement as incumbent-protected SCC neighborhood search (IPSNS). This protection gives a simple invariant: the final solution cannot be worse than the best internal seed under the same objective and nonnegative-weight assumptions. The invariant is not an approximation guarantee, but it prevents refinement from degrading the seed and makes the local search behavior auditable instance by instance.

The computational study is organized around five evidence streams. The main sparse benchmark evaluates public weighted directed graph instances derived from graph-benchmark collections, with negative-weight instances separated from the standard nonnegative comparisons. An ablation study isolates the effect of the topological add-back phase, the use of the best seed, the incumbent-protected refinement, and the SCC-priority choice on a representative subset. An exact small-instance validation compares heuristic orderings against a bitmask dynamic-programming optimum where exact optimization is feasible. A sparse external-baseline study compares the proposed framework with the open-source DRMacIver/FAS heuristic, the Eades heuristic as implemented in python-igraph, a weighted adaptation of Eades et al. (1993), a Borda net-score ordering, and random multi-start. Finally, a dense transfer study evaluates whether the framework carries over to complete weighted linear-ordering instances from the Linear Ordering Problem Library (LOLIB). This combination was chosen to address earlier

concerns about novelty framing, topological extraction dependence, parameter justification, baseline breadth, dataset scope, and venue fit.

The results support a deliberately bounded claim. On 105 sparse benchmark instances, the incumbent-protected refinement has no incumbent violations relative to either internal seed. In the ablation study, the add-back phase reduces mean backward weight by 5.9 percent relative to local-ratio without add-back, and the refinement gives an additional 0.75 percent mean reduction on the tested subset. On 57 standard small nonnegative instances, the refinement is optimal on 56 instances, with mean relative gap about 0.0006 percent. On 97 standard nonnegative sparse instances in the external comparison, it achieves the best observed backward weight on 96 instances; the strongest external baseline, DRMacIver, is about 21.6 percent worse in mean backward weight on this sparse benchmark. These numbers support strong performance on sparse nonnegative weighted digraphs, but they do not support a universal dominance claim.

The dense LOLIB study is therefore reported as a scope boundary rather than hidden as an unfavorable side result. LOLIB instances are complete weighted ordering instances, closely related to the linear ordering problem and structurally different from the sparse directed graphs that motivate the framework. On 50 dense LOLIB instances, the tournament-native DRMacIver/FAS heuristic obtains the best observed solution on 45 instances, while the proposed refinement is best on 5 instances and has about 3.88 percent higher mean backward weight. This outcome is consistent with the structural distinction: SCC-local refinement is useful when sparse cyclic subgraphs provide meaningful neighborhoods, whereas dense complete tournaments reward algorithms specialized to pairwise-comparison structure. The manuscript therefore positions the method as a reproducible general weighted-digraph framework with strong sparse performance and an explicitly documented dense-ordering limitation.

The contributions of this paper are as follows:

- We present an engineered local-ratio topological add-back heuristic for nonnegative weighted feedback arc set instances, separating inherited local-ratio ideas from the implementation and add-back choices studied here.
- We introduce incumbent-protected SCC neighborhood refinement, a local-search framework that preserves non-worsening relative to the best internal seed and can be audited for incumbent violations.

- We provide a multi-part computational evaluation covering sparse public benchmarks, component ablation, exact small-instance validation, external sparse baselines, and dense LOLIB transfer testing.
- We state explicit claim boundaries for negative-weight exclusions, dense linear-ordering instances, weighted Eades adaptation, reproducibility, and the absence of a new approximation guarantee.

The remainder of the paper is organized as follows. Section 2 reviews feedback arc set heuristics, local-ratio methods, ordering heuristics, exact methods, and linear-ordering benchmarks. Section 3 defines the minimum weighted feedback arc set objective, the ordering-induced backward weight, and the relation to dense linear-ordering notation. Section 4 describes the local-ratio topological add-back procedure, the weighted seed, and the incumbent-protected SCC refinement. Section 5 presents datasets, baselines, metrics, and reproducibility controls. Section 6 summarizes the computational results. Section 7 discusses limitations, including the dense LOLIB boundary, and Section 8 concludes.

2. Related work

2.1. Feedback arc set and local-ratio methods

The feedback arc set problem and its weighted variants sit at the intersection of graph optimization, ordering, and ranking. In directed graphs, the problem can be phrased either as deleting a minimum-weight set of arcs to obtain an acyclic graph or as finding an ordering with minimum backward weight. That equivalence makes the problem relevant both to graph-theoretic cycle breaking and to applications that begin from pairwise precedence or preference data. The approximation literature is broader for vertex-deletion or undirected variants than for the weighted directed setting considered here, but it still provides important methodological context. In particular, Bar-Yehuda et al. Bar-Yehuda et al. (1998) place local-ratio reasoning in a general approximation framework, showing how weighted combinatorial objectives can be decomposed while preserving feasible improvement structure. Their setting is not the same as minimum weighted feedback arc set on directed graphs, and we do not rely on their theorem to obtain a new approximation guarantee here. However, the paper is directly relevant to the design logic of iterative weight reductions on conflict structures.

For directed feedback problems, Demetrescu and Finocchi Demetrescu and Finocchi (2003) provide the closest methodological antecedent to the cycle-reduction component used in this manuscript. Their contribution is important for two reasons. First, it treats directed feedback problems algorithmically rather than only as ordering formulations. Second, it shows how local-ratio reductions can be applied to directed cycles in a way that is compatible with practical heuristics. The local-ratio step in the present work is inherited in spirit from that line of research: when a directed cycle is found, all arcs on the cycle are reduced by the minimum weight on the cycle, at least one arc becomes tight, and the active graph is simplified. What is new here is not the local-ratio principle itself, but an engineered sparse-graph implementation that records removed arcs explicitly, reconstructs a deterministic topological order once the active graph becomes acyclic, and then runs a heavy-first add-back step to reduce unnecessary deletions.

This distinction matters for positioning. It would be inaccurate to claim that the manuscript introduces a new local-ratio algorithm for feedback arc set in the theoretical sense. The contribution is better understood as an algorithm-engineering framework built around a known reduction principle. The emphasis is on deterministic data handling, auditable behavior, compatibility with sparse weighted digraphs, and integration with a second-stage refinement process. That framing is more defensible for a computational journal such as *Computers & Industrial Engineering*, and it also addresses a recurring novelty concern in prior manuscript iterations: the paper must be clear about what is inherited and what is newly engineered.

2.2. Heuristics and software baselines

Heuristic work on feedback arc set and related ordering tasks spans greedy graph rules, local search, and software packages that expose ordering or arc-deletion primitives. A classical baseline is the greedy heuristic of Eades, Lin, and Smyth Eades et al. (1993). Their method is attractive because it is fast, interpretable, and structurally simple: it repeatedly extracts sources, sinks, or high net-score vertices to construct an order. Even when a modern study ultimately outperforms such a heuristic, it remains an essential reference point because many later practical systems either implement that rule directly or build on the same score-based intuition. In the present paper, Eades-style ideas appear in two places: as the historical baseline that motivates fast greedy ordering, and as the conceptual starting point for external and adapted software baselines used in the experiments.

The external-baseline design in this manuscript intentionally mixes library-grade and specialized open-source tools. The python-igraph documentation for `Graph.feedback_arc_set` python-igraph developers (2026) exposes both an Eades-style heuristic and exact integer-programming modes. It is a useful baseline because it reflects what a practitioner could call from a general-purpose graph library. In contrast, the DRMacIver/FAS tool DRMacIver (2026) is a specialized weighted-ordering method rather than a general graph-analysis package. Its importance here is empirical: it is the strongest external baseline on the sparse standard benchmark and the dominant method on the dense LOLIB transfer test.

The remaining baselines in the experiments are deliberately simple and interpretable. Weighted net-score ordering and random multistart are not intended to represent state of the art; they provide calibration points. The weighted Eades adaptation used here should also be interpreted carefully. The original Eades–Lin–Smyth paper Eades et al. (1993) does not define the weighted adaptation used in this study, so the manuscript describes that baseline as a local adaptation rather than attributing the weighted form to the original authors. That distinction is part of the methodological hygiene of the study: external baselines should be credited as external when they are external, and local adaptations should be labeled as such even when they are inspired by well-known methods.

There is a broader family of learning-based ranking methods that could be discussed without entering the experimental comparison set. GNNRank He et al. (2022), for example, frames ranking from pairwise comparisons as a trainable graph-neural ranking problem. It is relevant as a signal that the ranking literature has moved beyond purely combinatorial heuristics, but it is not used as a baseline here because the present manuscript targets deterministic, training-free graph heuristics rather than learned ranking models.

2.3. Exact methods and ordering formulations

The exact side of the literature serves two different roles in this paper. First, it sets expectations about computational hardness. Second, it provides the reference standard for the small-instance validation experiment. Baharev et al. Baharev et al. (2021) give a modern exact treatment of minimum feedback arc set through formulations that connect the problem to broader combinatorial optimization machinery. Their work is not the computational baseline for the full sparse benchmark in this manuscript; the sparse benchmark is too large for exact optimization to be the main comparison mode.

Instead, exact methods enter the study through a limited validation regime on small instances, where the question is not whether exact optimization scales globally, but whether the heuristic framework is near-optimal when exact answers are available.

This distinction also helps separate feedback arc set from the larger family of ordering models. Dense linear ordering formulations are often written as forward-weight maximization on complete matrices, whereas sparse directed graph formulations arise more naturally as backward-weight minimization on incomplete digraphs. The two are mathematically linked, but they are not computationally identical in practice.

The present manuscript deliberately moves between the two views. The sparse benchmark is the main claim-bearing environment and is presented in graph terms. The LOLIB transfer test is presented in linear-ordering terms because those instances are dense complete ordering problems. Keeping both formulations visible is important. Otherwise the reader could mistakenly treat good sparse-graph performance as evidence of dense linear-ordering dominance, which the experiments do not support.

At larger scales, exact computation is often replaced by decomposition, approximation, or specialized engineering. Simpson et al. Simpson et al. (2016) provide one useful point of reference by treating feedback arc set computation in a web-scale setting. Their work is not directly comparable to the exact small-instance validation in this manuscript, but it illustrates the broader computational reality: scalable feedback arc methods are usually driven by structure, heuristics, or deployment constraints rather than by globally exact optimization. This matters for the framing of the current paper. The goal is not to displace exact methods; it is to present a heuristic framework whose computational behavior is strong on the sparse benchmark and whose limitations are visible rather than obscured.

2.4. Linear ordering benchmarks and manuscript position

The dense-transfer side of the paper is best understood through the linear ordering literature. Marti, Reinelt, and Duarte Marti et al. (2012) provide a benchmark-oriented account of the Linear Ordering Problem Library (LOLIB) and its heuristic evaluation culture. That paper is the most appropriate formal reference for LOLIB-style benchmark usage in the current manuscript because it anchors the library in a published benchmark-comparison setting rather than only as a download page. The official LOLIB

page itself Marti (2010) remains important for provenance and dataset identification, especially because benchmark users often obtain the instances directly from library pages or their mirrors. Together, these sources explain why dense complete tournaments have a different methodological status from sparse DIMACS-style directed graphs: they are not merely larger graphs of the same type, but instances drawn from a different benchmark tradition with different algorithmic expectations.

The sparse benchmark side has a different provenance. The graph-benchmarks collection Ali Dasdan (2026) is a public set of directed instances in DIMACS-style format rather than a linear-ordering library.

That difference supports a central design choice: the sparse and dense benchmarks test different claims. The sparse benchmark is the main evidence for practical value on nonnegative weighted digraphs. The dense benchmark is a transfer test used to determine where a general weighted-digraph heuristic stops being competitive against dense-ordering-native methods.

The ranking-from-pairwise-comparisons viewpoint connects these benchmark families conceptually. In dense complete settings, every ordered pair carries evidence, so forward-weight maximization aligns naturally with tournament and ranking formulations. In sparse settings, many pairwise relations are absent, heterogeneous, or structurally constrained by applications such as dependency repair or precedence resolution.

This is exactly why a method that performs well on sparse graph-benchmarks may still lose on LOLIB. The issue is not inconsistency in the experiments; it is a change in problem structure. Related work therefore has to prepare the reader for a comparative message that is strong but bounded.

This paper occupies a specific position in the literature. It is not a new approximation-theory paper in the style of local-ratio foundations Bar-Yehuda et al. (1998) or directed-feedback algorithm design Demetrescu and Finocchi (2003). It is not an exact-method paper in the style of Baharev et al. Baharev et al. (2021). It is not a web-scale systems paper in the style of Simpson et al. Simpson et al. (2016). It is also not a dense linear-ordering or trainable ranking paper in the style of the LOLIB benchmark tradition Marti et al. (2012) or graph-neural ranking work such as GNNRank He et al. (2022). Instead, it is an integrated computational heuristic study whose main object is a reproducible framework for nonnegative weighted directed graphs.

That framework combines three roles that are often separated in the literature: a cycle-reduction seed grounded in prior local-ratio ideas, a complementary weighted seed, and a refinement phase that acts only when it can

preserve or improve the incumbent objective. The practical significance of that combination is twofold. First, it gives a transparent explanation for why the method should be robust on sparse directed graphs: the initial seeds construct feasible orders in different ways, and the refinement concentrates effort on strongly connected regions where ordering uncertainty remains. Second, it yields a claim that is strong enough to matter and narrow enough to defend: the framework is highly effective on the main sparse benchmark, near-optimal on small exact-validation instances, and explicitly limited on dense LOLIB instances where the external tournament-native baseline is stronger.

For a journal-paper contribution, that position is appropriate. The manuscript does not need to claim universal dominance or theoretical novelty where the evidence is not there. What it can claim is that a carefully engineered local-ratio-and-refinement framework materially improves reproducible performance on sparse weighted directed ordering problems. It also remains explicit about its dense-ordering limits.

3. Problem definition

Let $G = (V, A, w)$ be a directed graph with vertex set V , arc set $A \subseteq V \times V$, and nonnegative arc weights $w : A \rightarrow \mathbb{R}_{\geq 0}$. A feedback arc set is a set $F \subseteq A$ such that removing F makes the graph acyclic. The minimum weighted feedback arc set problem asks for a minimum-weight feedback arc set, i.e., a set F minimizing $\sum_{a \in F} w(a)$.

An ordering π assigns each vertex a distinct position. An arc (u, v) is backward under π if u appears after v . The ordering-induced backward weight is

$$\text{bw}(\pi) = \sum_{(u,v) \in A: \pi(u) > \pi(v)} w(u, v). \quad (1)$$

For any ordering π , the set of backward arcs is a feedback arc set: all remaining forward arcs point from earlier to later positions and therefore form an acyclic graph. Conversely, every acyclic subgraph admits a topological order. The weighted feedback arc set problem can therefore be treated as the problem of finding an ordering with minimum backward weight.

Dense linear ordering problem benchmarks are often stated in terms of a complete weight matrix C , where an ordering $\pi = (\pi_1, \dots, \pi_n)$ is evaluated by the forward weight. In that notation,

Table 1: Algorithm components and their manuscript roles.

Component	Primary mechanism	Role in this paper
LR-TA	Global local-ratio cycle reduction followed by heavy-first add-back	Engineered reproducible seed; main constructive method derived from prior local-ratio ideas
WMSF-style seed	Ordered arc removal, minimization, stabilization, and remimization inside SCCs	Complementary internal seed and baseline; broadens incumbent quality before refinement
IPSNS	SCC-local destroy-repair moves with restricted local-ratio repair and restricted add-back	Main refinement framework; improves only when the global objective decreases

$$\text{fw}_C(\pi) = \sum_{1 \leq i < j \leq n} C_{\pi_i, \pi_j}, \quad \text{bw}_C(\pi) = \sum_{1 \leq i < j \leq n} C_{\pi_j, \pi_i} = W_C - \text{fw}_C(\pi), \quad (2)$$

where $W_C = \sum_{i \neq j} C_{ij}$ is the total off-diagonal weight. Thus maximizing forward weight on a complete dense linear-ordering instance is equivalent to minimizing backward weight after converting the matrix to a complete weighted digraph. The manuscript reports LOLIB results as a transfer test. Those instances are dense complete ordering problems rather than sparse directed graph benchmarks.

4. Proposed methodology

The framework combines three layers: a local-ratio topological add-back seed, a complementary weighted seed, and an incumbent-protected refinement over strongly connected components. Figure 1 summarizes the data flow, while Tables 1 and 2 isolate the roles of each component and the invariants that the implementation preserves.

Table 2: Implementation-level invariants emphasized in the manuscript.

Property	Meaning
Acyclic seed output	LR-TA and the WMSF-style seed both output an acyclic active graph that induces a valid ordering.
Heavy-first add-back	Reinserted edges are processed in nonincreasing original weight, so high-weight arcs are given priority for restoration.
Deterministic topological order	Kahn-style topological reconstruction uses deterministic tie-breaking, making seed output reproducible for fixed input.
Incumbent protection	IPSNS updates the incumbent only when the candidate backward weight is strictly smaller.
Reproducible randomization	IPSNS SCC selection and destroy fractions are fully determined by the configured random seed.

The implementation reads an aggregated weighted DIMACS digraph and stores it in edge-indexed arrays. Each arc receives an edge identifier together with fixed endpoint arrays and an activity flag. This design is shared across LR-TA, WMSF-style seeding, and IPSNS. The common objective is the backward weight defined in Section 3; all moves are evaluated under that objective, and all determinism claims in this section refer to fixed input graphs and fixed random seeds.

4.1. Local-ratio topological add-back

LR-TA is the first constructive component. It inherits the idea of local-ratio cycle reduction from the directed-feedback setting of Demetrescu and Finocchi Demetrescu and Finocchi (2003), but the implementation studied here is an engineered sparse-graph instantiation rather than a new theorem. The algorithm maintains an active subgraph over the original edge identifiers and repeatedly applies two phases.

In Phase I, the algorithm searches the active graph for any directed cycle using an iterative depth-first search. The search is deterministic because vertices and outgoing edge lists are scanned in fixed order, and only the first detected cycle is used. If C is the detected cycle, LR-TA computes

$$\varepsilon(C) = \min_{e \in C} w(e) \quad (3)$$

and reduces the mutable weight of every arc in C by $\varepsilon(C)$. At least one arc becomes tight. Any arc whose reduced weight falls below the numerical tolerance is deactivated and appended to the removed-edge set F_{LR} . The phase terminates when no directed cycle remains in the active graph, so the active subgraph is a directed acyclic graph.

In Phase II, LR-TA attempts to reinsert removed arcs without recreating cyclicity. Removed arcs are processed in nonincreasing order of original weight, with deterministic tie-breaking on endpoints and edge identifier. Let π be a topological order of the current active directed acyclic graph and let $\text{rank}_\pi(v)$ denote the position of vertex v . A candidate arc (u, v) is accepted immediately if

$$\text{rank}_\pi(u) < \text{rank}_\pi(v), \quad (4)$$

because the current topological order already certifies that the insertion is forward. Otherwise the algorithm checks whether v can reach u in the current active graph, using a reachability search pruned by the topological rank interval. The arc is reinserted if and only if no such path exists. Whenever a backward candidate is accepted, the topological order is recomputed so that future admissibility tests remain exact. The extracted topological order is deterministic under the implementation used here, but it is not mathematically unique: different linear extensions of the same DAG can exist. The paper therefore treats the chosen extraction rule as an engineering convention rather than as a canonical ordering of an acyclic state.

Algorithm 1 summarizes the implementation.

The add-back phase is where this implementation departs most clearly from a bare local-ratio reduction. It is also the component with the strongest ablation signal in the current experiments. The manuscript should therefore treat LR-TA as a two-phase engineered seed rather than as a one-step local-ratio routine.

4.2. WMSF-style seed construction

The second seed is a weighted removeArcs–minimize–stabilize pipeline implemented locally as an internal baseline. The role of this component in the manuscript is pragmatic: it broadens the seed space available to the refinement stage and helps define the incumbent-protection guarantee. It should not be presented as the main novelty of the paper.

At a high level, the WMSF-style seed decomposes the graph into strongly connected components and processes each nontrivial component separately.

Algorithm 1 LR-TA seed construction

Require: Weighted digraph $G = (V, A, w)$ with $w \geq 0$

Ensure: Ordering π_{LR} , removed-edge set F_{LR}

```
1: Build edge-indexed arrays  $(U, V, W_0, W, \text{active}, \text{adj})$ 
2:  $F_{\text{LR}} \leftarrow \emptyset$ 
3: while the active graph contains a directed cycle  $C$  do
4:    $\varepsilon \leftarrow \min_{e \in C} W[e]$ 
5:   for all  $e \in C$  do
6:      $W[e] \leftarrow W[e] - \varepsilon$ 
7:     if  $W[e] \leq \tau$  and  $e$  is active then
8:       deactivate  $e$ ; set  $W[e] \leftarrow 0$ ; add  $e$  to  $F_{\text{LR}}$ 
9:     end if
10:  end for
11: end while
12: Compute deterministic topological order  $\pi$  of the active graph
13: for all  $e = (u, v) \in F_{\text{LR}}$  in nonincreasing order of  $W_0[e]$  do
14:   if  $\text{rank}_\pi(u) < \text{rank}_\pi(v)$  then
15:     reactivate  $e$ ; remove  $e$  from  $F_{\text{LR}}$ 
16:   else if  $v \not\rightarrow u$  in the current active graph then
17:     reactivate  $e$ ; remove  $e$  from  $F_{\text{LR}}$ 
18:     recompute deterministic topological order  $\pi$ 
19:   end if
20: end for
21: return  $\pi_{\text{LR}}, F_{\text{LR}}$ 
```

Within a component, it first removes arcs according to a prescribed weight-based ordering. Two orderings are supported in code: *L1*, which sorts arcs primarily by weight, and *L2*, which normalizes arc weight by a weighted in/out-degree denominator. The resulting removed set is then minimized by heavy-first add-back, followed by a stabilization step and a second minimization pass. When the entire graph is a single nontrivial strongly connected component, both *L1* and *L2* are tried and the lower-backward-weight seed is kept. Otherwise the configured ordering is used componentwise.

The point of carrying this seed into IPSNS is not to establish a separate literature claim about WMSF. Instead, it provides a complementary structural bias to LR-TA. LR-TA is driven by cycle reductions in the active graph. The WMSF-style seed is driven by ordered arc removals and subsequent min-

imization. Because the two constructions fail differently on some instances, taking the better of them before refinement produces a stronger and more robust incumbent than either seed alone.

4.3. Incumbent-protected SCC-neighborhood search

IPSNS begins by computing both seeds on the full graph: the full WMSF-style seed and the LR-TA seed. Their backward weights are evaluated under the common ordering objective, and the better one becomes the initial incumbent. The refinement stage then iterates over strongly connected components of the original aggregated graph, but the selection is filtered by the incumbent ordering: only components with positive current backward contribution are eligible, and a weighted-random choice is drawn from the top- K contributors by SCC backward weight. This makes the neighborhood search concentrated rather than exhaustive.

Within the chosen component, IPSNS applies a destroy-and-repair move. The destroy stage has two parts. First, a fraction of currently removed internal edges is reactivated in heavy-first order, temporarily giving the neighborhood more freedom. Second, a small fraction of currently active internal edges is forcibly removed in light-first order, encouraging escape from a locally stable order. The repair stage then applies a restricted local-ratio cycle reduction on that SCC neighborhood and finishes with restricted heavy-first add-back minimization. All graph edits are confined to the chosen component’s internal edges, while the rest of the incumbent graph remains untouched during the move.

If the repaired candidate fails to admit a valid topological order, the entire SCC-local state is reverted. If the candidate is valid, the full incumbent objective is recomputed. The incumbent is updated only if the new backward weight is strictly smaller than the current best value. Otherwise the move is rejected and the previous incumbent is restored exactly. This is what makes the method incumbent-protected: refinement never commits a worsening move.

Algorithm 2 summarizes the IPSNS driver at manuscript level.

Two implementation details are especially important for interpretation. First, the default WMSF seed mode inside IPSNS is the full per-SCC pipeline, not a weaker legacy approximation; this is required so that the refinement never starts from a seed weaker than the standalone WMSF baseline. Second, the randomization in IPSNS is controlled entirely by a fixed seed. The resulting sequence of SCC choices and destroy fractions is therefore reproducible

given the same input and parameter settings. The parameter choices themselves are interpreted conservatively later in the paper: the ablation study supports them as reasonable defaults on the tested subset, not as universally optimal settings.

4.4. Algorithmic invariants

The most important invariant in the framework belongs to IPSNS rather than to LR-TA. Let $\pi^{(0)}$ denote the better of the two internal seeds and let $\pi^{(t)}$ denote the incumbent ordering after t accepted-or-rejected IPSNS iterations. Because the incumbent is replaced only by a strictly improving candidate, the sequence satisfies

$$\text{bw}(\pi^{(t+1)}) \leq \text{bw}(\pi^{(t)}) \leq \text{bw}(\pi^{(0)}) \quad \text{for all } t \geq 0. \quad (5)$$

This monotone non-degradation property is not an approximation guarantee, but it is a concrete algorithmic promise: the final IPSNS solution is never worse than the better internal seed under the same objective and nonnegative-weight assumptions.

LR-TA and the WMSF-style seed have their own feasibility invariants. LR-TA always exits its reduction stage with an acyclic active graph, and its add-back phase only reactivates arcs that preserve acyclicity. The WMSF-style seed similarly alternates removal and minimization steps while checking acyclic feasibility. IPSNS inherits those feasibility properties locally during SCC repair and adds the stronger incumbent-protection rule globally. That combination explains why the paper can make a narrow but defensible guarantee about refinement quality without claiming any new theoretical approximation ratio.

5. Experimental design

The experimental design follows a deliberately narrow claim strategy. The manuscript does not attempt to demonstrate that one heuristic dominates every weighted feedback arc set or linear ordering regime. Instead, it evaluates a reproducible framework for nonnegative weighted directed graphs, using sparse benchmark instances as the primary setting and dense complete ordering instances as a transfer test. This organization responds directly to the main concerns identified in the earlier audit: baseline fairness, dataset breadth, exact-validation wording, parameter justification, and clear scope

Table 3: Overview of the five computational components used in the manuscript.

Study	Purpose	Inst.	Main setting
Internal benchmark	Safety check for LR-TA, WMSF, and IPSNS	105	sparse weighted DIMACS graphs
Ablation study	Add-back, refinement, and iteration study	10	representative sparse subset
Exact validation	Comparison with exact bitmask dynamic programming	57	standard nonnegative instances with $n \leq 20$
Sparse comparison	External baselines on the main benchmark	97	standard nonnegative sparse instances
Dense transfer test	Scope-boundary analysis on LOLIB instances	50	LOLIB complete weighted ordering instances

boundaries matter more here than simply increasing the number of experiments.

Five empirical questions motivate the design. The first is whether the full refinement framework improves on its internal seeds without violating the non-degradation invariant. The second is which components materially contribute to the final solution quality. The third is how close the framework comes to exact optimization on instances where exact optimization is feasible. The fourth is how the framework compares with strong internal and external baselines on the main sparse benchmark. The fifth is how far the sparse-benchmark conclusions transfer to dense complete ordering instances. Table 3 summarizes how the five computational components are used to answer these questions.

5.1. Datasets

The study uses two public benchmark families with different roles in the argument.

The primary benchmark family comes from the public **graph-benchmarks** collection Ali Dasdan (2026). After path-level deduplication, this yields 105 unique sparse weighted DIMACS digraph instances. These instances support the internal benchmark, the ablation study, the exact small-instance validation, and the external sparse comparison.

The standard analysis is restricted to the nonnegative-weight subset. Eight instances are excluded from the standard sparse comparison set because they contain negative-weight arcs: **gerez**, **howard-max**, **k3_3**, **ku**, **peterson**, **peterson1**, **peterson2**, and **stg0**. This exclusion is methodological rather than cosmetic. The local-ratio construction and the interpretation of incumbent protection in this manuscript are framed for nonnegative weights, so mixing negative-weight instances into the main tables would weaken the meaning of the guarantee. After exclusions, the primary sparse comparison set contains 97 standard instances.

The exact small-instance study uses the subset of this sparse benchmark for which exact optimization is feasible by the dynamic-programming implementation used in the study. The verified summaries show 57 standard nonnegative instances with $n \leq 20$. This subset is not treated as a separate benchmark family; it is a validation slice used to judge heuristic quality against exact optima.

The dense transfer test uses a 50-instance subset of the Linear Ordering Problem Library (LOLIB) 2010 Marti et al. (2012); Marti (2010). These instances differ structurally from the sparse benchmark in a way that matters algorithmically. They are complete weighted ordering instances, and the selected subset spans three families: SGB (25 instances), IO (10 instances), and RandA1 (15 instances across $n = 100, 150$, and 200). In the converted DIMACS representation used in the experiments, every ordered pair carries weight, so these instances are effectively dense complete tournaments or dense linear-ordering inputs rather than sparse general digraphs.

This distinction is central to the manuscript’s scope. The sparse benchmark is the main source of positive claims. LOLIB is not treated as coequal evidence for the primary contribution; it is used to determine where a sparse weighted-digraph heuristic remains competitive and where a tournament-native ordering method becomes preferable. Taken together, the sparse benchmark and LOLIB broaden the empirical surface enough to support a CAIE submission while keeping the comparisons interpretable.

5.2. Algorithms and baselines

The evaluation includes the proposed refinement framework, its internal seed methods, an exact reference method for the small-instance study, and a set of external or simple baseline heuristics.

The full method of interest is IPSNS, which starts from the better of two internal seeds and applies incumbent-protected SCC-local refinement.

The two seeds are LR-TA and WMSF. LR-TA is the engineered local-ratio plus topological add-back construction described in Section 4; WMSF is a weighted removeArcs–minimize–stabilize seed that is carried forward as a baseline and seed source rather than as a novel contribution of the present paper.

For the exact small-instance validation, the reference method is exact bitmask dynamic programming. That solver is not used as a scalable comparator on the full sparse benchmark; its role is only to certify optimality on the small validation subset.

The sparse and dense comparison studies also include five additional baselines:

- the DRMacIver/FAS tool DRMacIver (2026), treated as an external wrapper around an open-source specialized algorithm;
- python-igraph’s Eades-style feedback arc interface python-igraph developers (2026);
- a weighted Eades adaptation implemented locally and described explicitly as an adaptation rather than as the original Eades et al. (1993) algorithm;
- Borda net-score ordering;
- random multistart ranking.

This baseline mix is deliberate. It does not attempt to exhaust every recent ranking model or every exact solver. In particular, it does not include a learned model such as GNNRank, nor does it include a commercial exact solver for the full sparse dataset. Those omissions are stated explicitly. The intended comparison surface is the ecosystem of deterministic or near-deterministic weighted-digraph heuristics and accessible external tools that a practitioner could reasonably invoke without building a separate learning or solver pipeline.

Several fairness rules govern the comparisons. First, all quality comparisons use the same backward-weight objective on the original nonnegative arc weights. Second, official external methods and local adaptations are distinguished explicitly. The weighted Eades method is an adaptation of the unweighted Eades idea Eades et al. (1993); it is not presented as an official external implementation. Third, DRMacIver/FAS is reported exactly as it

was run in the study. The sparse comparison therefore retains its four incompletions on the standard set, and the dense comparison retains its strong overall performance. Finally, outcomes are interpreted in scope: losing to a tournament-native method on dense LOLIB is not treated as evidence against the sparse-graph claim, and strong sparse performance is not used to imply dense linear-ordering dominance.

5.3. *Evaluation metrics*

The primary objective throughout the paper is backward weight (BW), defined in Section 3. Lower BW is always better, and all main tables are organized around that quantity.

Several secondary metrics support interpretation. The number of global-best instances records how often a method attains the minimum observed BW among the compared algorithms. For the external sparse comparison, relative excess over IPSNS records the mean percentage by which a method’s BW exceeds IPSNS on completed instances. For the exact small-instance study, the mean gap from exact optimum and the number of optimal instances summarize how close each heuristic comes to exact optimization. Mean runtime is reported for practical context, although it is not normalized across machines. The dense complete-ordering study also reports forward ratio as a secondary descriptive measure of how much pairwise weight is oriented forward by the returned ranking. Finally, the incumbent-violation count records how often IPSNS would return a solution worse than LR-TA or WMSF; under the intended invariant this count should remain zero on the nonnegative setting.

These metrics correspond directly to the audit concerns. BW and global-best counts address baseline strength. Exact gap addresses the request for stronger quality validation without overstating theory. Runtime provides an engineering interpretation for CAIE readers. The incumbent-violation count ties the experiments back to the framework’s main formal promise. Forward-ratio and family-level dense results help articulate the sparse-versus-dense scope boundary rather than leaving it implicit.

5.4. *Implementation and reproducibility*

The manuscript relies only on recorded experimental outputs and generation scripts. No new experiment is run for the present writing pass. Instead, paper assets are regenerated from the saved summaries by a paper-specific build script. That script reads the verified JSON, CSV, and Markdown summaries for the five computational components, checks the key manuscript

numbers against their source files, regenerates the paper tables, and produces the result figures as vector PDF files. It also writes machine-readable and human-readable provenance notes for auditability.

This generation layer serves two purposes. First, it reduces manual transcription risk. Second, it makes the manuscript tables auditable: the paper-level assets are derived from saved summaries rather than being hand-maintained copies with unclear provenance.

The experimental runs themselves use fixed seeds where randomness exists and consistent wrappers for the compared methods. For IPSNS, the saved summaries record the default full WMSF seed mode and the iteration budgets used in each study. The ablation study is especially important here because it justifies those defaults conservatively: on the 10-instance subset, the mean BW for 50 iterations is identical to the 200-iteration mean, and the no-SCC-priority variant is nearly identical to the default. The manuscript interprets those results as subset-level support for reasonable defaults, not as universal parameter optimality.

The reproducibility story also includes transparent handling of exclusions and incompletions. Negative-weight sparse instances are excluded from the standard claim-bearing benchmark because they do not match the nonnegative setting studied in the manuscript. DRMacIver/FAS incompletions on the sparse benchmark are reported explicitly rather than hidden in filtered tables. The dense LOLIB transfer test is retained even though it narrows the paper’s claim surface, because that honesty is part of the contribution: the framework is strong on sparse weighted digraphs and bounded on dense complete ordering instances.

One remaining detail is administrative rather than algorithmic. The manuscript is prepared under double-anonymized review constraints, so code-availability language remains anonymized at manuscript level. That packaging issue does not affect the experimental evidence summarized here, but it does shape how reproducibility is described in the review version.

6. Computational results

The computational evidence should be read in the same order as the experimental design. The paper first establishes the main sparse-benchmark result, then checks exact small-instance quality, then uses ablations to interpret the engineering choices, and finally tests transfer to dense LOLIB instances. This ordering matters because the positive claim is intentionally

Table 4: Comparison on the 97 standard nonnegative sparse instances used for the primary claim-bearing benchmark. BW denotes backward weight; lower values are better. Relative excess is the mean percentage by which an algorithm’s BW exceeds IPSNS on completed instances.

Method	Complete	Mean BW	Mean RT (s)	Best	Relative excess
IPSNS (ours)	97/97	37,698	21.92	96	0.00%
LR-TA (ours)	97/97	38,327	0.08	80	0.71%
WMSF seed	97/97	40,005	1.31	61	2.06%
DRMacIver/FAS	93/97	53,173	4.00	56	21.61%
igraph Eades	97/97	95,920	0.01	40	30.54%
Weighted Eades	97/97	99,689	0.11	42	30.48%
Borda net score	97/97	512,277	0.00	27	55.55%
Random multistart	97/97	1,075,258	0.03	42	49.76%

DRMacIver/FAS is the strongest external comparator on this benchmark, but it completes 93 of 97 instances because two DAG cases return empty-tournament errors and two large sparse instances time out in the wrapper used here.

narrow: IPSNS is presented as a strong, reproducible heuristic framework for nonnegative sparse weighted directed graphs, not as a universal best method for every ordering benchmark.

6.1. Sparse weighted graph benchmarks

The primary result comes from the standard nonnegative sparse benchmark, with the full internal benchmark supplying the safety check for the refinement stage. Table 4 summarizes the main comparison on 97 instances, and Figures 2 and 3 visualize the gap structure and the distribution of instance wins.

The main outcome is clear on this benchmark. IPSNS achieves the lowest observed backward weight on 96 of the 97 standard instances. Its mean BW is 37,698, compared with 38,327 for LR-TA and 40,005 for WMSF. Among the external methods, DRMacIver/FAS is the strongest comparator, but its mean BW is still 53,173 on the completed standard instances, or about 21.61% above IPSNS. The remaining baselines are substantially weaker on average, with the Eades variants about 30% higher than IPSNS and the simpler Borda and random baselines far behind.

This result addresses the baseline-strength concern directly. The sparse-benchmark claim is not based only on a comparison among internal variants. It survives contact with a strong external comparator, a widely available

graph-library heuristic, and several simple but informative ranking baselines. The most important external comparison is DRMacIver/FAS, not because it wins the benchmark, but because it is the only external method that remains genuinely competitive on a large portion of the sparse set. Even there, the gap is small only on one instance: `r20_60`, where DRMacIver/FAS beats IPSNS by 3 BW units. That same instance is also the only IPSNS near-miss in the exact-validation study below, which makes the sparse story internally consistent rather than contradictory.

The full internal benchmark strengthens this interpretation by showing that the refinement stage does not obtain these gains by occasionally damaging its seeds. Across 105 sparse benchmark instances, IPSNS is never worse than LR-TA and never worse than WMSF, with zero incumbent-protection violations and zero nonempty errors in the verified summary. It strictly improves over LR-TA on 16 instances and over WMSF on 36 instances. The mean runtime per instance rises from 0.074 seconds for LR-TA and 1.24 seconds for WMSF to about 20.2 seconds for IPSNS, so the gain is not free; however, it is gained without giving up the best internal seed. That property is more informative than a generic claim that a hybrid beats its components because it connects the empirical behavior back to the monotone non-degradation invariant in Section 4.

The runtime pattern also clarifies the engineering tradeoff. LR-TA is the fastest strong method in the study: its mean sparse-benchmark runtime is only 0.08 seconds per instance, yet its mean BW is within about 1.7% of IPSNS. That makes LR-TA a credible low-latency choice when refinement time is limited. IPSNS then occupies the higher-quality point on the tradeoff curve, paying more runtime to secure the best observed solutions on almost every instance. For a CAIE framing, that is the practical interpretation that matters: the paper provides a fast seed heuristic, a stronger refinement framework, and a clear statement of what the extra computation buys.

The sparse results also address the concern that the paper should say more about dataset breadth without pretending to cover every graph family. The benchmark is not tiny, synthetic-only, or internal-only. It contains 97 standard nonnegative sparse instances after exclusions, and the comparison includes eight algorithms in total. At the same time, the claim remains intentionally scoped to this setting. The paper does not generalize from this table to negative-weight cases or to dense complete ordering benchmarks.

Table 5: Exact validation on the 57 standard nonnegative instances with $n \leq 20$. Mean gap is measured against exact bitmask dynamic programming on each instance.

Method	Optimal	Optimality rate	Mean gap
IPSNS (ours)	56/57	98.2%	0.0006%
LR-TA (ours)	55/57	96.5%	0.0590%
WMSF seed	51/57	89.5%	0.0961%

This table supports empirical near-optimality on small instances only; it does not imply a general approximation guarantee. The only IPSNS near-miss is instance `r20_60`.

Table 6: Ablation study on 10 representative sparse instances. Relative change is measured against LR-TA. Negative percentages indicate improvement.

Variant	Mean BW	Relative change	Mean RT (s)
LR without add-back	4525.1	+5.94%	0.017
LR-TA	4271.5	+0.00%	0.047
WMSF seed	4332.5	+1.43%	0.040
Best seed, no refinement	4271.5	+0.00%	0.069
IPSNS, 50 iterations	4239.2	-0.76%	0.778
IPSNS, 200 iterations	4239.2	-0.76%	5.734
IPSNS, no SCC priority	4239.1	-0.76%	5.727

The add-back phase accounts for the largest mean improvement (about 5.9%), while IPSNS contributes a further 0.75% on this subset. The equal 50- and 200-iteration means support the default iteration budget only as subset-level evidence, not as a universal tuning rule.

6.2. Exact validation and ablation evidence

The exact small-instance validation and the ablation study serve different purposes, but together they explain why the framework is both credible and interpretable. Table 5 reports the exact validation results, and Table 6 reports the component ablation.

The exact small-instance study supports a quality-validation narrative rather than a theoretical one. On the 57 standard nonnegative instances with $n \leq 20$, IPSNS is exactly optimal on 56 and has a mean gap of 0.0006% from the exact optimum. LR-TA is optimal on 55 instances with mean gap 0.0590%, and WMSF is optimal on 51 with mean gap 0.0961%. These numbers show that the framework is not merely competitive in relative benchmark

terms; on small instances where exact optimization is feasible, it is essentially exact in practice.

The wording matters here. This table does not provide an approximation guarantee, a worst-case bound, or a proof of global behavior beyond the validated subset. The correct statement is that IPSNS is empirically near-optimal on the exact-validation subset. That statement is already strong enough for the role this study is meant to play.

The exact-validation result also helps interpret the single sparse-benchmark loss to DRMacIver/FAS. The only IPSNS near-miss is again `r20_60`. This is also the only standard small instance where IPSNS fails to match exact optimum. That consistency reduces the chance that the sparse comparison is driven by unstable reporting or artifact mismatch. Instead, it suggests a small, structurally difficult case where the framework is near optimum but not uniquely dominant.

The ablation study explains where the framework’s gains come from. The add-back phase is the dominant contributor among the tested design choices. Removing it raises mean BW from 4,271.5 for LR-TA to 4,525.1 for the no-add-back variant, a degradation of about 5.94%. This is the clearest experimental support for treating topological add-back as a substantive engineering contribution rather than a decorative post-processing step.

The refinement stage then provides a smaller but still measurable gain. Moving from LR-TA to IPSNS with 50 iterations reduces mean BW from 4,271.5 to 4,239.2, about 0.76%. In absolute terms that gain is much smaller than the add-back gain, but it is still meaningful because it is obtained after starting from a strong seed and while preserving the non-degradation invariant. This is the kind of marginal improvement one expects from a bounded refinement layer: it does not re-invent the seed; it improves the hard residue that remains after seeding.

The ablation results also provide the manuscript’s main evidence for parameter justification, and the interpretation must remain conservative. On this 10-instance subset, the 50-iteration and 200-iteration IPSNS variants have the same mean BW, which suggests that the default budget is already beyond the point of visible average improvement on these instances. Likewise, removing SCC priority changes the mean only from 4,239.2 to 4,239.1, effectively negligible at the reported precision. These observations support the default settings as reasonable choices, but they do not show that those settings are universally optimal on other graph families or scales.

Table 7: Dense LOLIB transfer test on 50 complete weighted ordering instances. Lower BW is better. The upper panel reports overall performance; the lower panel isolates the two most relevant methods by family to make the dense-scope boundary explicit.

Method	Mean BW	Best	Mean RT (s)
DRMacIver/FAS	571,687	45/50	1.01
IPSNS (ours)	582,354	5/50	37.92
LR-TA (ours)	582,948	2/50	0.22
WMSF seed	585,926	1/50	0.19
igraph Eades	659,570	0/50	0.01
Weighted Eades	675,235	0/50	0.01
Random multistart	883,664	0/50	0.02
Borda net score	1,066,045	0/50	0.00
<i>Family breakdown for IPSNS and DRMacIver/FAS</i>			
Family	IPSNS best	DRMacIver/FAS best	Mean BW (IPSNS / DRMacIver/FAS)
SGB	1/25	24/25	1,048,056 / 1,036,085
IO	4/10	6/10	36,996 / 32,860
RandA1	0/15	15/15	169,755 / 156,909

DRMacIver/FAS is tournament-native and is the best overall dense LOLIB method in this study (45/50 instances, 3.88% lower mean BW than IPSNS). This dense transfer study is therefore a scope boundary, not evidence for universal dominance on dense linear-ordering problems.

6.3. Dense LOLIB transfer test

Table 7 and Figure 4 summarize the dense LOLIB transfer test.

This is the result that prevents overclaiming. On the 50 dense complete ordering instances, DRMacIver/FAS is the best overall method in the study: it wins 45 instances and achieves mean BW 571,687, compared with 582,354 for IPSNS. The mean gap is 3.88% in DRMacIver/FAS’s favor. IPSNS remains second overall and retains zero incumbent violations against LR-TA and WMSF, but the dense benchmark does not support any claim of overall dominance by the proposed framework.

The family breakdown explains why. DRMacIver/FAS wins 24 of the 25 SGB instances and all 15 RandA1 instances. IPSNS is more competitive on IO, where it wins 4 of 10 against DRMacIver/FAS’s 6 of 10, but that does not overturn the overall dense-ordering result. The cleanest interpretation is structural: DRMacIver/FAS is tournament-native, whereas IPSNS is designed around sparse weighted-digraph seeding and SCC-local refinement. When every ordered pair carries weight and the instance is fully dense, the advantage shifts toward methods specialized for that structure.

This finding matters because it narrows the manuscript’s claim in a cred-

ible way. A paper that reported only the sparse benchmark could be challenged for hiding unfavorable regimes. By retaining the dense transfer study and writing it honestly, the manuscript does the opposite. It shows that the framework transfers nontrivially to a dense benchmark, remains competitive on a structured IO subset, and still preserves the internal no-worsening invariant, but it also shows that dense complete ordering is not the framework’s strongest regime.

Across the section as a whole, the empirical message is coherent. The strongest claim belongs to the sparse nonnegative benchmark, where IPSNS is the best observed method on 96 of 97 standard instances and never returns a result worse than the better internal seed. The exact small-instance validation shows that the quality story is not merely relative to heuristics, and the ablation study shows that the gains are not coming from an uninterpretable black box. The dense transfer study then marks the boundary of that success case. This combination allows the paper to be assertive without becoming misleading: the framework is a strong and reproducible method for the sparse weighted-digraph regime, and the computational evidence shows how far that statement can be pushed.

7. Discussion

The computational evidence supports a clear but bounded interpretation of the proposed framework. The strongest result is the sparse nonnegative benchmark, where IPSNS achieves the lowest observed backward weight on 96 of 97 standard instances and never degrades the better internal seed. The importance of that result is not only that the framework performs well, but that it performs well in a way that matches the intended algorithmic design. The methodology combines two complementary global constructions with a refinement stage that concentrates effort on the cyclic parts of the graph and accepts a move only when the global objective improves. The sparse results therefore do not look accidental: they are consistent with the structure that the method is designed to exploit.

The sparse benchmark interpretation begins with the role of the seed methods. LR-TA and WMSF produce feasible orderings by different mechanisms. LR-TA is driven by local-ratio reductions and topological add-back, whereas WMSF follows a removal-and-minimization template. On many sparse instances those two constructions are already strong, and LR-TA in particular offers an excellent quality–runtime tradeoff. IPSNS improves on

that base by focusing its neighborhood moves on strongly connected components that still contribute substantial backward weight under the incumbent ordering. This is a natural strategy for sparse cyclic graphs. When the graph contains large acyclic regions and a smaller cyclic core, indiscriminate perturbation is wasteful; the search budget is better spent where ordering ambiguity remains concentrated. The empirical advantage of IPSNS on the sparse benchmark is therefore not simply that it runs longer than LR-TA. It is that the extra computation is directed at the right structural target.

The internal safety result is central to this interpretation. On the full 105-instance internal benchmark, IPSNS never returns a solution worse than LR-TA or WMSF. That is more than a convenient side condition. In practical decision-support or ordering pipelines, a refinement stage that occasionally damages a good seed is hard to trust, hard to tune, and difficult to diagnose. The monotone non-degradation rule makes the framework auditable. If refinement finds no helpful move, the best seed is preserved. If refinement improves the incumbent, that improvement is directly visible in the final objective. For an engineering-oriented journal audience, this kind of operational reliability is often more useful than a loosely justified claim that a local search method is “usually better.”

The exact small-instance validation strengthens the quality story without changing its theoretical status. IPSNS is exact on 56 of 57 standard nonnegative instances in the validation subset and has a mean gap of 0.0006%. That is strong evidence that the heuristic framework is solving the small instances almost perfectly in practice. It is not, however, evidence for a new approximation ratio or a new optimality guarantee beyond the validated subset. That distinction matters because heuristic papers are often tempted to over-read small exact studies. Here the correct interpretation is narrower and stronger: when exact optimization is feasible, the framework tracks the optimum closely; when the instances become larger, the paper relies on comparative benchmark evidence rather than on a theoretical claim it does not possess.

The ablation study helps clarify which parts of the framework are carrying the result. The topological add-back phase is the dominant design element among the tested components, with a mean backward-weight reduction of about 5.9% relative to local-ratio without add-back. That is the main reason LR-TA should be understood as an engineered instantiation rather than as a thin wrapper around prior local-ratio ideas. The refinement stage then contributes a smaller additional improvement, about 0.75% on the ablation

subset. This is a plausible pattern rather than a disappointing one. Seed construction determines most of the ordering quality; refinement cleans up the remaining difficult cases without sacrificing the incumbent. The observed near-equality between 50 and 200 IPSNS iterations on the ablation subset also clarifies how the parameter choices should be discussed. They appear reasonable and stable on the tested subset, but they should not be promoted as universally optimal defaults for every future graph family.

The dense LOLIB transfer test is the key scope boundary. On that benchmark, DRMacIver/FAS is stronger overall, winning 45 of 50 instances and outperforming IPSNS in mean backward weight. This result should not be treated as a side note or a caveat hidden in fine print. It is part of the scientific contribution because it shows where the framework stops being the most competitive option among the methods tested. The structural explanation is straightforward. DRMacIver/FAS is tournament-native, and LOLIB supplies dense complete ordering instances where every ordered pair contributes information. IPSNS, by contrast, is built around sparse weighted-digraph seeding and SCC-local refinement. When the graph is fully dense, the distinction between cyclic core and acyclic background becomes much less informative, and the advantage shifts toward methods specialized for dense pairwise-comparison structure.

At the same time, the dense result should not be overinterpreted in the opposite direction. IPSNS remains second overall on LOLIB, preserves its non-degradation property, and is competitive on the IO family, where it wins 4 of 10 instances. The correct message is not that the method fails on dense complete ordering problems. The message is that dense complete ordering is not the paper’s primary target and that a tournament-native method is preferable there. This kind of honest scope boundary improves the paper rather than weakening it. A computational optimization article is more credible when it identifies where its methodology is strong, where it is merely competitive, and where a different structural assumption favors another class of methods.

Several limitations remain and should be stated directly. First, the paper does not offer a new approximation guarantee. The local-ratio principle used in LR-TA is prior art, and the present contribution is an engineered reproducible instantiation plus the incumbent-protected SCC refinement framework. Second, the main claims apply only to nonnegative-weight instances. Negative-weight sparse instances are excluded from the standard analysis because they do not match the intended setting of the framework and would

blur the meaning of the reported invariant. Third, the exact-validation study is limited to small instances where bitmask dynamic programming is feasible. It supports the quality narrative but does not remove the need for heuristic benchmarking on larger graphs. Fourth, dense linear-ordering instances are not the primary target, and the LOLIB results show that clearly. Fifth, the review version of the paper is anonymized, so data and code availability must be described in double-blind form until the submission process allows a public release.

These limitations also help identify the practical implications for a CAIE audience. The paper is best read as a contribution to weighted ordering and cyclic-dependency repair in graph-based decision-support settings. Many engineering and operations workflows contain precedence conflicts, pairwise preference information, or cyclic dependencies that must be resolved into an ordering without claiming that every arc or relation can be simultaneously respected. In such settings, a methodology that starts from strong fast seeds, improves only when the objective truly decreases, and remains reproducible under fixed seeds and explicit tie handling can be more valuable than a theoretically stronger method that is harder to deploy or interpret. The sparse benchmark results suggest that the framework is well suited to that style of application: graphs are large enough to make exact optimization impractical, sparse enough that SCC-local refinement remains meaningful, and structured enough that the combination of LR-TA, WMSF, and IPSNS improves over simpler ranking heuristics.

Future work should follow from those boundaries rather than ignore them. One direction is richer parameter tuning, especially to understand when longer refinement budgets or alternative SCC-selection rules materially improve sparse instances beyond the current defaults. Another is broader evaluation on ranking and application-facing datasets that better expose domain structure than generic benchmarks alone. A third is exact or ILP-based polishing on medium-sized instances if solver setup becomes available in a later study. That would not replace the present heuristic contribution, but it could sharpen the quality reference point between the small exact-validation regime and the large sparse benchmark regime. A fourth is broader study of negative-weight or mixed-sign settings under a formulation that preserves a clean interpretation of the objective and the refinement invariant.

Overall, the manuscript is strongest when it remains faithful to its empirical boundaries. It presents a reproducible methodology for the minimum weighted feedback arc set problem in nonnegative weighted directed graphs,

demonstrates clear strength on the main sparse benchmark, and documents where a different structural regime favors a different algorithmic family. That is a useful and defensible contribution for a computational optimization paper.

8. Conclusions

This paper studied the minimum weighted feedback arc set problem through its ordering interpretation on nonnegative weighted directed graphs. The main contribution is not a new approximation framework, but a reproducible computational methodology that combines an engineered local-ratio construction with topological add-back, a complementary weighted seed, and an incumbent-protected SCC-local refinement stage. Within that design, LR-TA serves as a strong constructive seed, while IPSNS is the main refinement framework that improves the incumbent only when the global backward-weight objective decreases.

The computational evidence supports a clear scope-bounded conclusion. On the standard nonnegative sparse benchmark, IPSNS attains the lowest observed backward weight on 96 of 97 instances and improves on all tested internal and external comparators, including the strongest external baseline. On the full internal benchmark, the refinement stage never worsens the better internal seed, which confirms the intended non-degradation behavior in practice.

The exact small-instance validation shows that the framework is near-optimal when exact optimization is feasible. The ablation study shows that the topological add-back phase is the dominant design element, with SCC-local refinement providing an additional but smaller gain.

The results also define an important boundary. On dense complete LOLIB ordering instances, the tournament-native DRMacIver/FAS baseline is stronger overall. This does not negate the sparse-benchmark contribution; it clarifies it. The proposed framework is strongest in the sparse nonnegative weighted-digraph regime for which it was designed. It should be interpreted as a general weighted-digraph heuristic rather than as a dense-native linear-ordering solver.

For a computational optimization audience, the practical value of the framework lies in this combination of quality, structural focus, and reproducibility. The methodology is suitable for weighted ordering and cyclic-dependency repair settings where exact optimization is too expensive, sparse

graph structure is meaningful, and refinement should not be allowed to damage a good seed. The manuscript therefore contributes a usable and auditable heuristic approach rather than a broader, less defensible universal claim. Within those boundaries, the results suggest that engineered local-ratio seeding plus incumbent-protected SCC refinement is an effective strategy for sparse weighted directed ordering problems.

Two immediate directions follow from this evidence. The first is broader evaluation on additional sparse weighted digraph families, especially application-derived instances with richer edge-weight semantics. The second is methodological: the same incumbent-protected refinement discipline could be combined with stronger dense-ordering seeds or exact large-neighborhood subroutines without changing the basic non-degradation philosophy that motivated the present framework.

CRediT authorship contribution statement

Author details are omitted from the review manuscript for double-anonymized submission. A complete CRedit statement will be provided in the separate title page and submission metadata.

Funding statement

This research did not receive external funding.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have influenced the work reported in this paper.

Data and code availability

An anonymized reproducibility artifact containing the implementation, experiment scripts, selected instance lists, conversion utilities, and postprocessed result tables is provided as supplementary material for review. The benchmark instances used in the study are publicly available from the sources cited in the manuscript; the artifact includes scripts and metadata needed to obtain and convert them. A public repository and archival DOI will be released upon acceptance.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used AI-assisted tools, including ChatGPT, Codex, Claude, and Perplexity AI, to support literature exploration, organization of experimental material, and language editing. The authors reviewed and edited all generated content, verified the relevant sources and experimental results, and take full responsibility for the content of the submitted manuscript.

References

- Ali Dasdan, 2026. graph-benchmarks: Directed graph benchmarks in dimacs graph format. <https://github.com/alidasdan/graph-benchmarks>. Dataset repository; accessed 2026-06-06.
- Baharev, A., Schichl, H., Neumaier, A., Achterberg, T., 2021. An exact method for the minimum feedback arc set problem. *ACM Journal of Experimental Algorithmics* 26, 1.4:1–1.4:28. doi:10.1145/3446429.
- Bar-Yehuda, R., Geiger, D., Naor, J.S., Roth, R.M., 1998. Approximation algorithms for the feedback vertex set problem with applications to constraint satisfaction and bayesian inference. *SIAM Journal on Computing* 27, 942–959. doi:10.1137/S0097539796305109.
- Demetrescu, C., Finocchi, I., 2003. Combinatorial algorithms for feedback problems in directed graphs. *Information Processing Letters* 86, 129–136. doi:10.1016/S0020-0190(02)00491-X.
- DRMacIver, 2026. Feedback-arc-set. <https://github.com/DRMacIver/Feedback-Arc-Set>. Software repository; experiment commit 16ff24a92fde886e58819180a9fe686e60991c5c; accessed 2026-06-06.
- Eades, P., Lin, X., Smyth, W.F., 1993. A fast and effective heuristic for the feedback arc set problem. *Information Processing Letters* 47, 319–323. doi:10.1016/0020-0190(93)90079-0.
- He, Y., Gan, Q., Wipf, D., Reinert, G.D., Yan, J., Cucuringu, M., 2022. Gn-rank: Learning global rankings from pairwise comparisons via directed

- graph neural networks, in: Proceedings of the 39th International Conference on Machine Learning, pp. 8581–8612. URL: <https://proceedings.mlr.press/v162/he22b.html>.
- Marti, R., 2010. Lolib: Linear ordering problem library. <https://www.uv.es/~rmarti/paper/lop.html>. Official library page; accessed 2026-06-06.
- Marti, R., Reinelt, G., Duarte, A., 2012. A benchmark library and a comparison of heuristic methods for the linear ordering problem. Computational Optimization and Applications 51, 1297–1317. doi:10.1007/s10589-010-9384-9.
- python-igraph developers, 2026. python-igraph documentation: Graph.feedback_arc_set. <https://python.igraph.org/en/stable/api/igraph.Graph.html>. Documentation for the weighted feedback arc set API; accessed 2026-06-06.
- Simpson, M., Srinivasan, V., Thomo, A., 2016. Efficient computation of feedback arc set at web-scale. Proceedings of the VLDB Endowment 10, 133–144. doi:10.14778/3021924.3021930.




Figure 1: Conceptual overview of the integrated framework. Two constructive seeds are evaluated first; IPSNS then applies incumbent-protected SCC-local destroy-and-repair moves to the better seed only when they improve the global objective.

Algorithm 2 IPSNS with incumbent protection

Require: Weighted digraph $G = (V, A, w)$, iteration budget T

Ensure: Ordering π^* with objective no worse than the best internal seed

```
1: Build WMSF-style seed  $\pi_W$  and LR-TA seed  $\pi_{LR}$ 
2:  $\pi^* \leftarrow \arg \min\{\text{bw}(\pi_W), \text{bw}(\pi_{LR})\}$ 
3: for  $t = 1, \dots, T$  do
4:   Score each nontrivial SCC by its current backward contribution
5:   if all SCC scores are zero then
6:     break
7:   end if
8:   Sample one SCC from the top- $K$  positive-score SCCs using weighted
   random choice
9:   Save the current internal SCC state
10:  Reactivate a heavy fraction of removed internal edges
11:  Remove a light fraction of active internal edges
12:  Apply SCC-restricted local-ratio repair
13:  Apply SCC-restricted heavy-first add-back minimization
14:  if the candidate SCC state is invalid then
15:    restore saved SCC state; continue
16:  end if
17:  Recompute the global ordering and candidate objective
18:  if the candidate objective is smaller than  $\text{bw}(\pi^*)$  then
19:    accept the candidate and update  $\pi^*$ 
20:  else
21:    restore the previous incumbent exactly
22:  end if
23: end for
24: return  $\pi^*$ 
```




Figure 2: Mean backward-weight excess over IPSNS on the 97 standard nonnegative sparse instances. Lower is better; IPSNS is the reference at zero.




Figure 3: Number of instances on which each method attains the lowest observed backward weight on the primary sparse benchmark.




Figure 4: Global-best counts by LOLIB family for IPSNS and DRMacIver/FAS. The split makes the dense-scope boundary explicit: DRMacIver/FAS is stronger overall, while IPSNS remains competitive on the IO family.